

UNIVERSITY OF MUMBAI
DEPARTMENT OF COMPUTER SCIENCE



मुंबई विद्यापीठ
University of Mumbai
Re-accredited with A++ Grade
(CGPA 3.65) by NAAC (3rd Cycle 2021)

M.Sc. Computer Science – Semester II
(NEP 2020)

Machine Learning

JOURNAL

2023-2024

Seat No. _____



मुंबई विद्यापीठ
University of Mumbai
Re-accredited with A++ Grade
(CGPA 3.65) by NAAC (3rd Cycle 2021)



UNIVERSITY OF MUMBAI
DEPARTMENT OF COMPUTER SCIENCE

CERTIFICATE

This is to certify that the work entered in this journal was done in the University
Department of Computer Science laboratory by
Mr./Ms. _____ Seat No. _____
for the course of M.Sc. (Computer Science) - Semester II (NEP 2020) during the
academic year 2023- 2024 in a satisfactory manner.

Subject In-charge

Head of Department

External Examiner

INDEX

Sr. no.	Name of the practical	Page No.	Date	Sign
1	Implement Linear Regression(Diabetes Dataset).			
2	Implement Logistic Regression. (Iris Dataset)			
3	Implement Multinomial Logistic Regression (Iris Dataset)			
4	Implement SVM Classifier (Iris Datasets)			
5	Train and fine-tune a Decision Tree for Moon dataset.			
6	Train an SVM regression on the California Housing Dataset.			
7	Implement Batch Gradient Descent with Early Stopping for Softmax Regression.			
8	Implement MLP for Classification of Handwritten digits(MNIST datasets)			
9	Classification of Image of clothing using Tensorflow(Fashion MNIST dataset)			
10	Implement Regression to predict fuel efficiency using TensorFlow (Auto MPG dataset).			

PRACTICAL-01

Aim: Implement Linear Regression(Diabetes Dataset).

Theory:

Linear regression analysis is used to predict the value of a variable based on the value of another variable. The variable you want to predict is called the dependent variable. The variable you are using to predict the other variable's value is called the independent variable.

Code :

```
import numpy as np
import pandas as pd
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns

# Load the Diabetes dataset from sklearn
diabetes = datasets.load_diabetes()

# Create a DataFrame from the dataset
diabetes_df = pd.DataFrame(data=diabetes.data,
columns=diabetes.feature_names)
diabetes_df['target'] = diabetes.target
```

Display basic information about the dataset(EDA)

```
print("Dataset Shape:", diabetes_df.shape)
```

```
print("\nColumns:", diabetes_df.columns)
```

```
print("\nInfo:")
```

```
print(diabetes_df.info())
```

```
print("\nNull Values:")
```

```
print(diabetes_df.isnull().sum())
```

```
Dataset Shape: (442, 11)
```

```
Columns: Index(['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6',
               'target'],
              dtype='object')
```

```
Info:
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 442 entries, 0 to 441
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null	Count	Dtype
0	age	442	non-null	float64
1	sex	442	non-null	float64
2	bmi	442	non-null	float64
3	bp	442	non-null	float64
4	s1	442	non-null	float64
5	s2	442	non-null	float64
6	s3	442	non-null	float64
7	s4	442	non-null	float64
8	s5	442	non-null	float64
9	s6	442	non-null	float64
10	target	442	non-null	float64

```
dtypes: float64(11)
```

```
memory usage: 38.1 KB
```

```
...
```

```
s5      0
```

```
s6      0
```

```
target  0
```

```
dtype: int64
```

Display summary statistics of numerical columns

```
print("\nSummary Statistics:")
```

```
print(diabetes_df.describe())
```

Summary Statistics:

	age	sex	bmi	bp	s1 \
count	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02
mean	-1.444295e-18	2.543215e-18	-2.255925e-16	-4.854086e-17	-1.428596e-17
std	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02
min	-1.072256e-01	-4.464164e-02	-9.027530e-02	-1.123988e-01	-1.267807e-01
25%	-3.729927e-02	-4.464164e-02	-3.422907e-02	-3.665608e-02	-3.424784e-02
50%	5.383060e-03	-4.464164e-02	-7.283766e-03	-5.670422e-03	-4.320866e-03
75%	3.807591e-02	5.068012e-02	3.124802e-02	3.564379e-02	2.835801e-02
max	1.107267e-01	5.068012e-02	1.705552e-01	1.320436e-01	1.539137e-01

	s2	s3	s4	s5	s6 \
count	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02
mean	3.898811e-17	-6.028360e-18	-1.788100e-17	9.243486e-17	1.351770e-17
std	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02
min	-1.156131e-01	-1.023071e-01	-7.639450e-02	-1.260971e-01	-1.377672e-01
25%	-3.035840e-02	-3.511716e-02	-3.949338e-02	-3.324559e-02	-3.317903e-02
50%	-3.819065e-03	-6.584468e-03	-2.592262e-03	-1.947171e-03	-1.077698e-03
75%	2.984439e-02	2.931150e-02	3.430886e-02	3.243232e-02	2.791705e-02
max	1.987880e-01	1.811791e-01	1.852344e-01	1.335973e-01	1.356118e-01

	target
count	442.000000
mean	152.133484

```
# Visualize the distribution of the target variable (diabetes progression)
```

```
plt.figure(figsize=(8, 6))
```


```
sns.histplot(diabetes_df['target'], bins=30, kde=True, color='blue')
```

```
plt.title('Distribution of Diabetes Progression')
```

```
plt.xlabel('Diabetes Progression')
```

```
plt.ylabel('Frequency')
```

```
plt.show()
```



Pairplot to visualize relationships between features and the target variable

```
plt.figure(figsize=(12, 10))
```

```
sns.pairplot(diabetes_df, diag_kind='kde')
```

```
plt.suptitle('Pairplot of Diabetes Dataset', y=1.02)
```

```
plt.show()
```

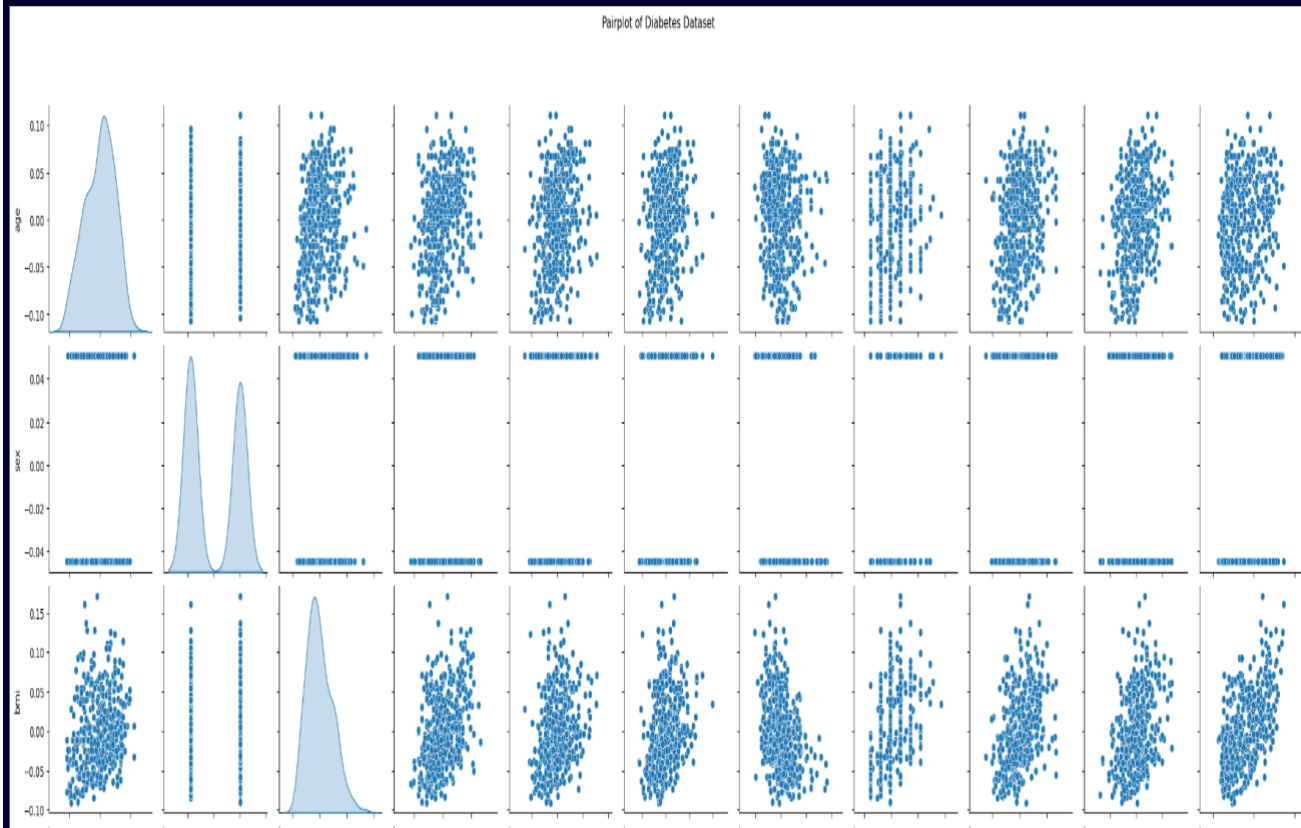
#Split the Data into Features and Target

Split the dataset into features (X) and target (y)

```
X = diabetes_df.drop('target', axis=1) # Features
```

```
y = diabetes_df['target'] # Target (continuous variable)
```

<Figure size 1200x1000 with 0 Axes>



```
# Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

```
# Initialize the Linear Regression model
```

```
model = LinearRegression()
```

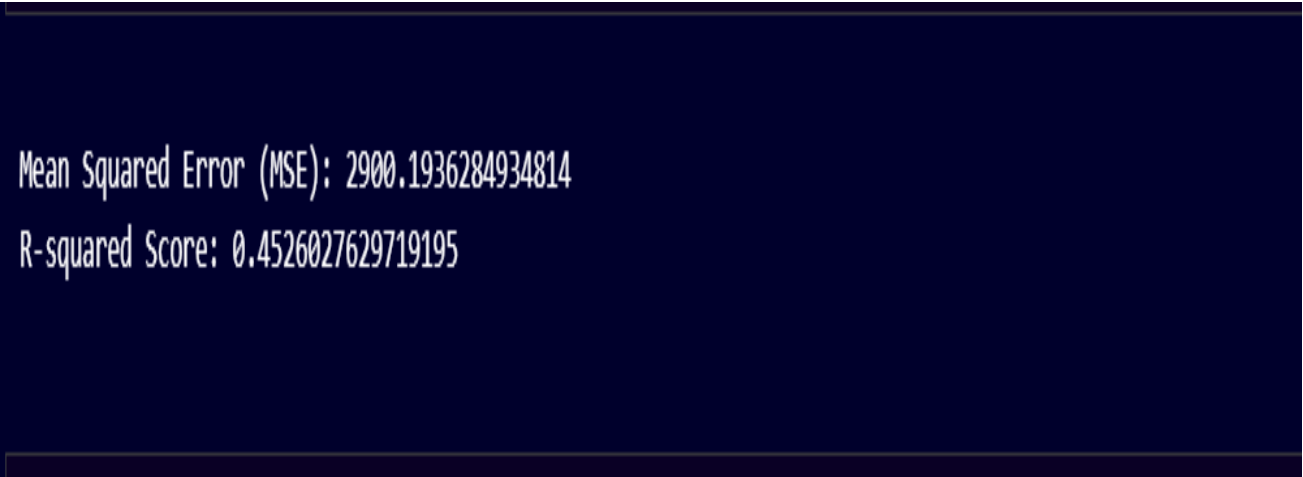
```
# Train the model on the training data
```

```
model.fit(X_train, y_train)
```



```
# Make predictions on the test data
y_pred = model.predict(X_test)

# Evaluate the model's performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("\nMean Squared Error (MSE):", mse)
print("R-squared Score:", r2)
```

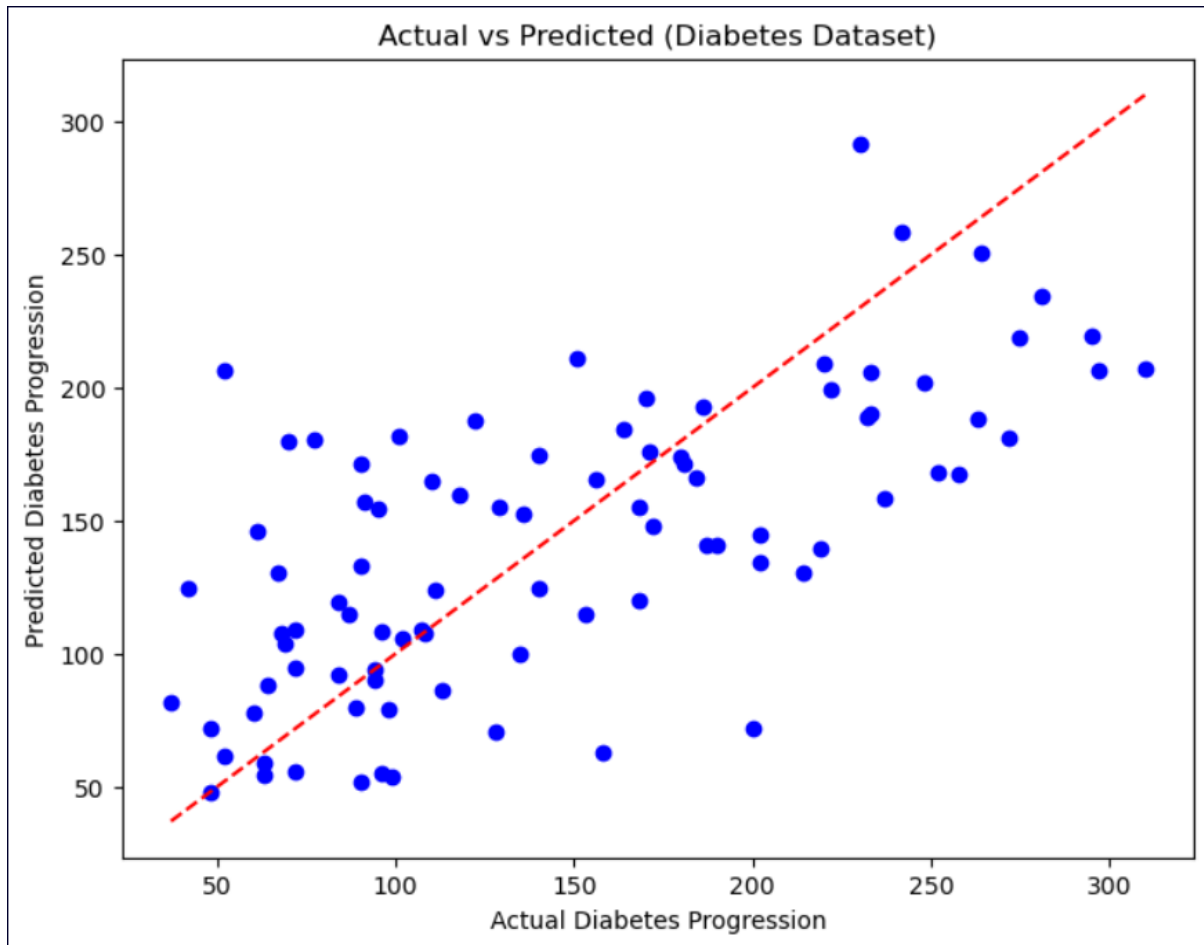


```
Mean Squared Error (MSE): 2900.1936284934814
R-squared Score: 0.4526027629719195
```

```
print("\nMean Squared Error (MSE):", mse)
print("R-squared Score:", r2)

# Plot predicted vs actual values
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], linestyle='-
-', color='red')
```

```
plt.xlabel('Actual Diabetes Progression')  
plt.ylabel('Predicted Diabetes Progression')  
plt.title('Actual vs Predicted (Diabetes Dataset)')  
plt.show()
```



Practical No 2

Aim: Implement Logistic Regression. (Iris Dataset)

Theory:

Logistic regression is a supervised machine learning algorithm widely used for binary classification tasks, such as identifying whether an email is spam or not and diagnosing diseases by assessing the presence or absence of specific conditions based on patient test results.

CODE:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

# Load the Iris dataset from sklearn
iris = datasets.load_iris()

# Convert the data into a DataFrame
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
iris_df['target'] = iris.target
```

`iris_df.shape`

```
(150, 5)
```

`iris_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sepal length (cm)      150 non-null   float64
1   sepal width (cm)       150 non-null   float64
2   petal length (cm)      150 non-null   float64
3   petal width (cm)       150 non-null   float64
4   target                 150 non-null   int32
dtypes: float64(4), int32(1)
memory usage: 5.4 KB
```

`iris_df.describe()`

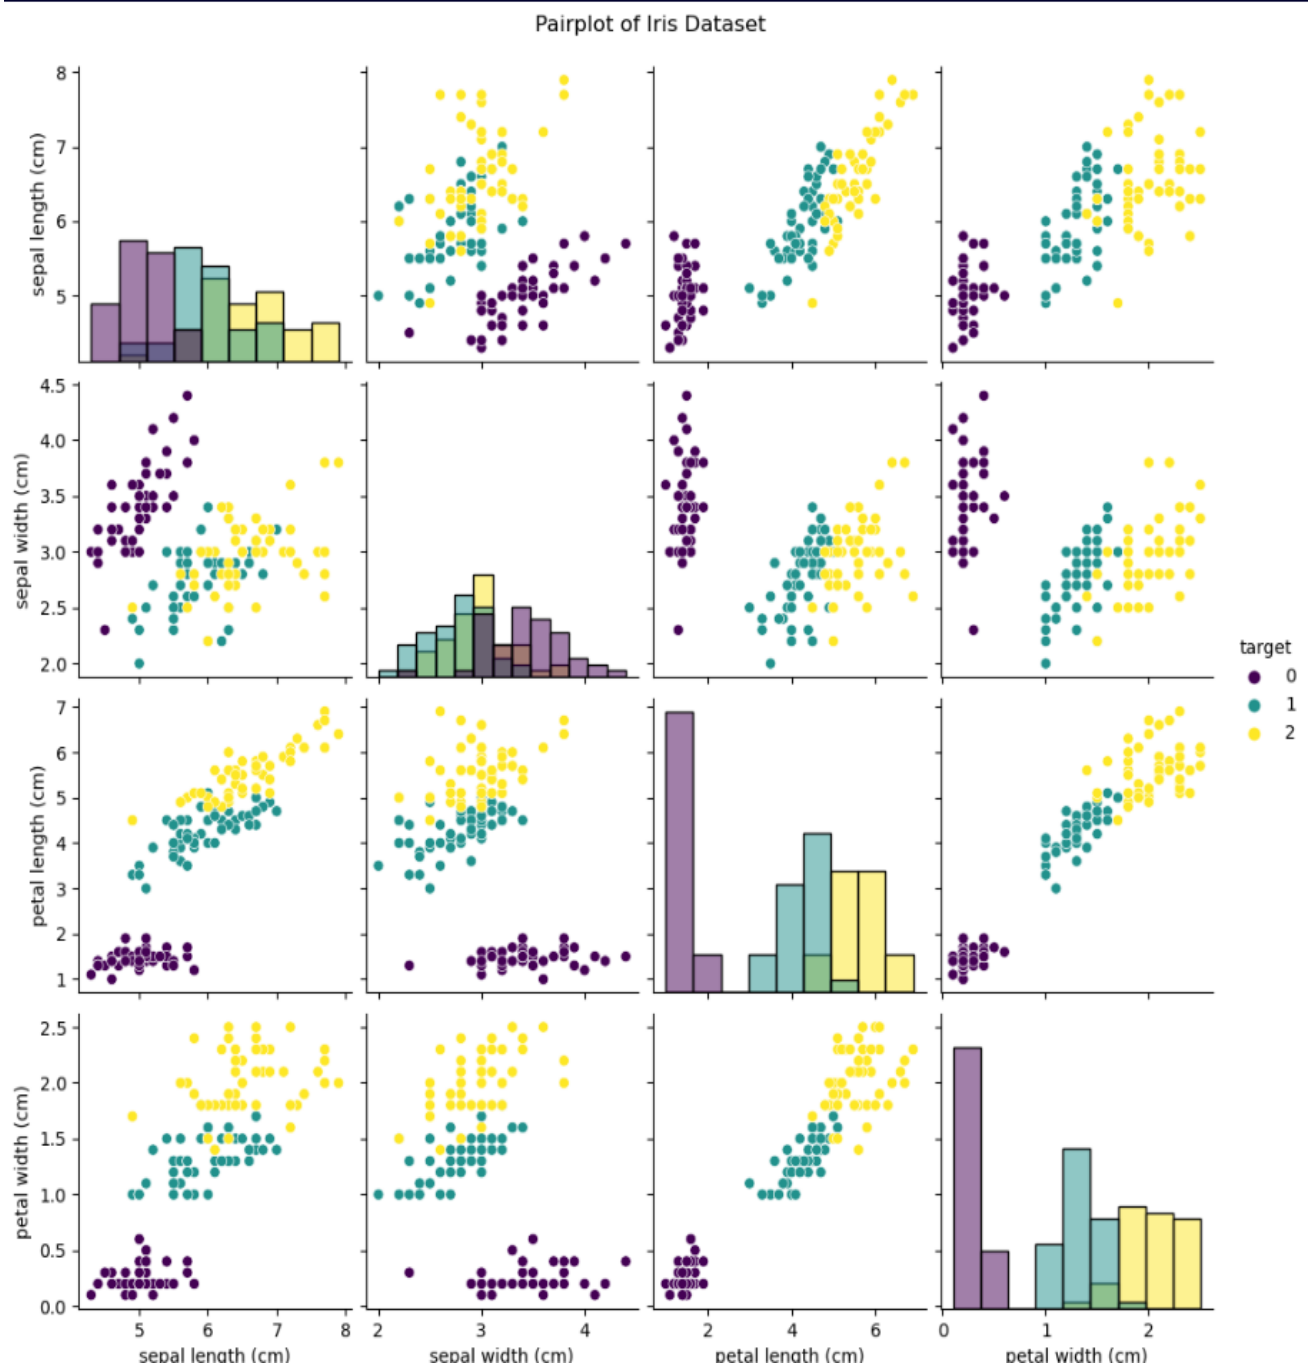
	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000
max	7.900000	4.400000	6.900000	2.500000	2.000000

EDA: Pairplot to visualize relationships between features

`sns.pairplot(iris_df, hue='target', palette='viridis', diag_kind='hist')`

`plt.suptitle("Pairplot of Iris Dataset", y=1.02)`

`plt.show()`



Other Plots

```
plt.figure(figsize=(12, 6))
```

Boxplot

```
plt.subplot(1, 2, 1)
```

```
sns.boxplot(x='target', y='sepal length (cm)', data=iris_df)
```

```
plt.title('Boxplot of Sepal Length by Target')
```

Violin Plot

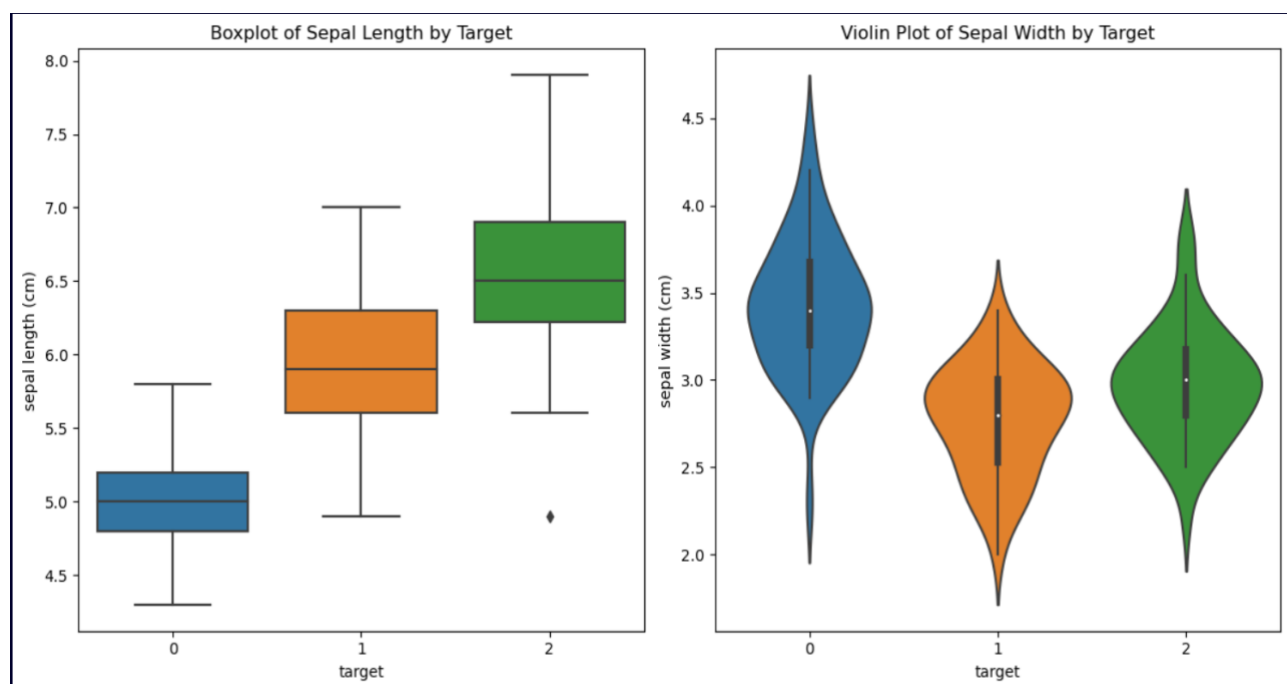
```
plt.subplot(1, 2, 2)
```

```
sns.violinplot(x='target', y='sepal width (cm)', data=iris_df)
```

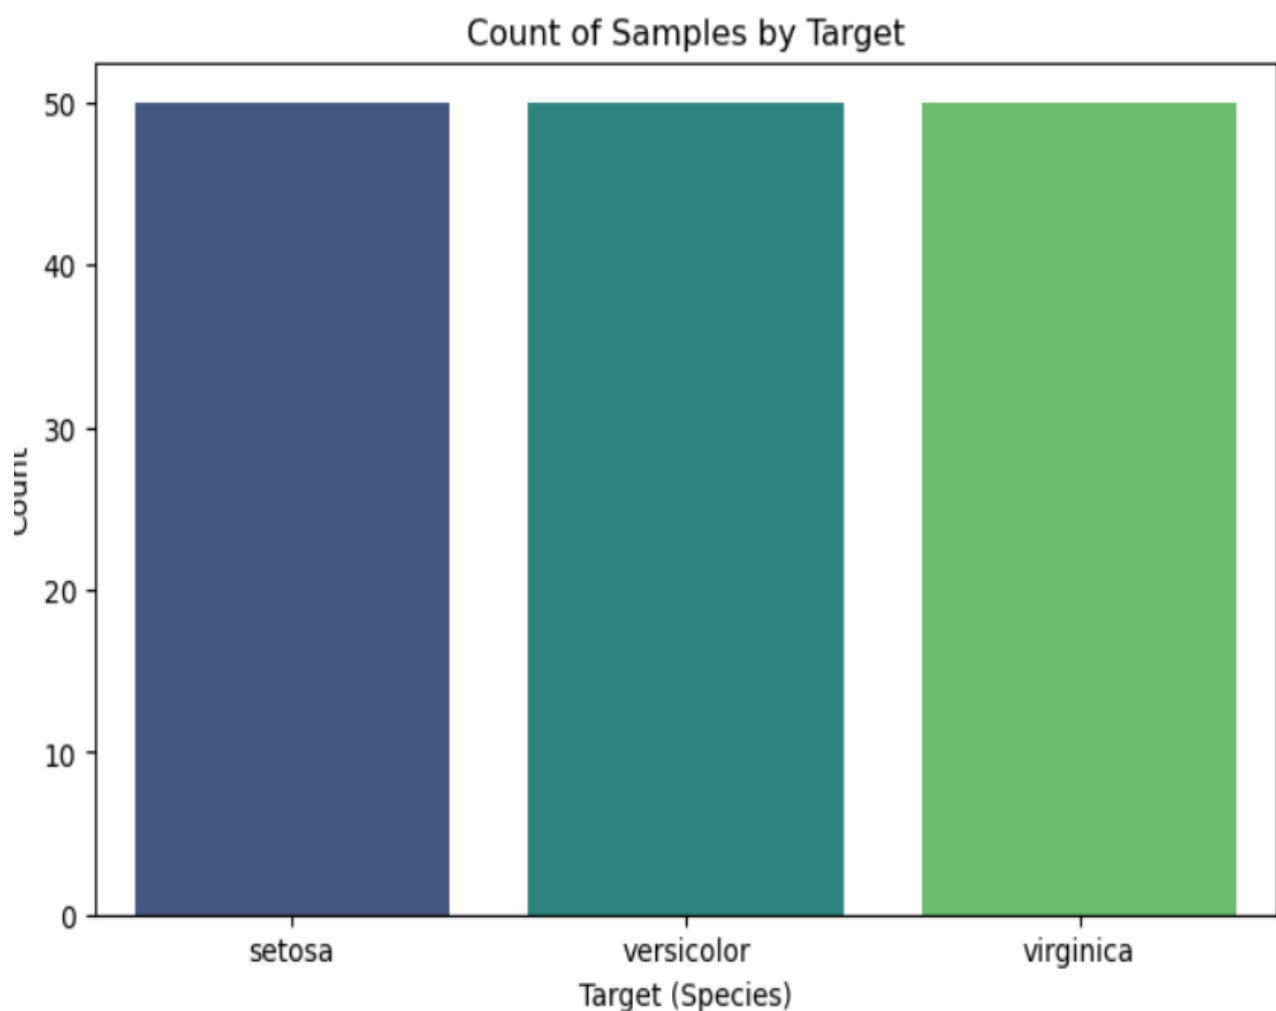
```
plt.title('Violin Plot of Sepal Width by Target')
```

```
plt.tight_layout()
```

```
plt.show()
```



```
# Count Plot (Bar Plot)
plt.figure(figsize=(8, 5))
sns.countplot(x='target', data=iris_df, palette='viridis')
plt.title('Count of Samples by Target')
plt.xlabel('Target (Species)')
plt.ylabel('Count')
plt.xticks(ticks=[0, 1, 2], labels=iris.target_names)
plt.show()
```



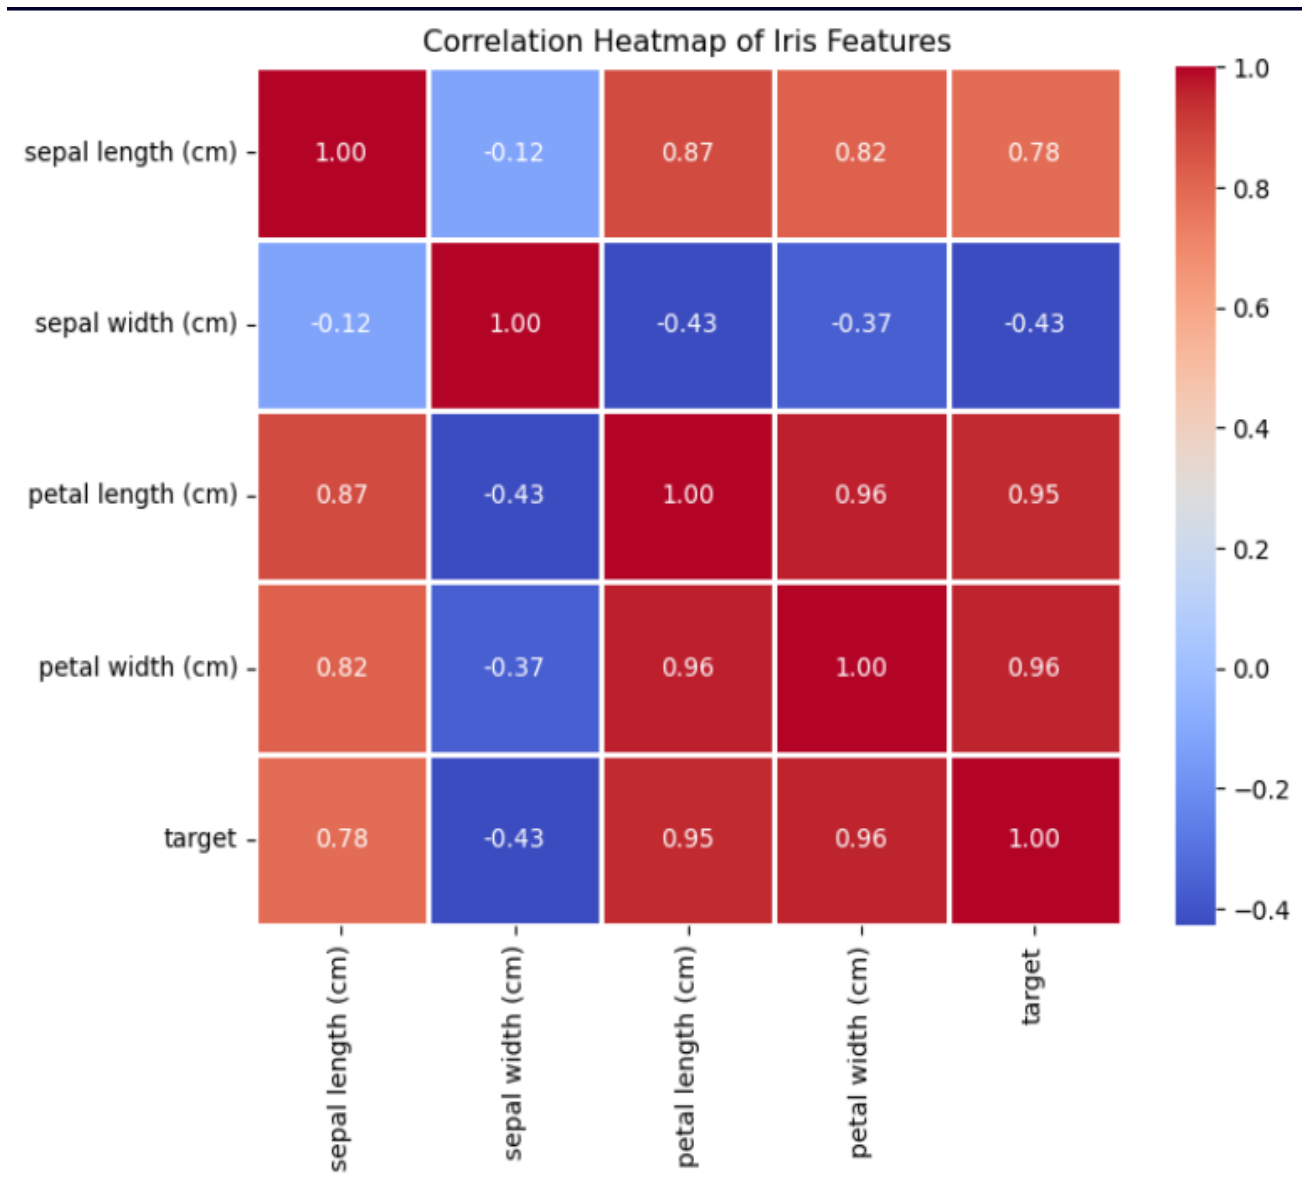
Correlation Heatmap

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(iris_df.corr(), annot=True, cmap='coolwarm', fmt='.2f',  
linewidths=1, linecolor='white')
```

```
plt.title('Correlation Heatmap of Iris Features')
```

```
plt.show()
```



Separate features (X) and target (y) from the DataFrame

X = iris_df.drop('target', axis=1) # Features

y = iris_df['target'] # Target (labels)

Split the dataset into training and testing sets

**X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)**

Standardize features by removing the mean and scaling to unit variance

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)

Initialize the Logistic Regression model

model = LogisticRegression()

Train the model on the training data

model.fit(X_train, y_train)

Predict on the test data

y_pred = model.predict(X_test)

Calculate training and testing accuracy

train_accuracy = accuracy_score(y_train, model.predict(X_train))

test_accuracy = accuracy_score(y_test, y_pred)

print("Training Accuracy:", train_accuracy)

print("Testing Accuracy:", test_accuracy)

```
Training Accuracy: 0.9666666666666667
Testing Accuracy: 1.0
```

Create a confusion matrix

conf_matrix = confusion_matrix(y_test, y_pred)

print("\nConfusion Matrix:")

print(conf_matrix)

```
Confusion Matrix:
```

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

```
# Print classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred,
target_names=iris.target_names))
```

```
Classification Report:
              precision    recall  f1-score   support

   setosa         1.00        1.00        1.00         10
  versicolor      1.00        1.00        1.00          9
   virginica      1.00        1.00        1.00         11

 accuracy         1.00        1.00        1.00         30
  macro avg       1.00        1.00        1.00         30
 weighted avg     1.00        1.00        1.00         30
```

Print probabilities of classification for the first few samples in the test set

```
print("\nProbabilities of Classification:")
probabilities = model.predict_proba(X_test[:5])
for i, prob in enumerate(probabilities):
    print(f"Sample {i+1}: {list(zip(iris.target_names, prob))}")
```

```
Probabilities of Classification:
Sample 1: [('setosa', 0.011457196118264068), ('versicolor', 0.8759785262009813), ('virginica', 0.11256427768075462)]
Sample 2: [('setosa', 0.9644113023352993), ('versicolor', 0.035588286377131434), ('virginica', 4.112875693141769e-07)]
Sample 3: [('setosa', 3.773229944007567e-08), ('versicolor', 0.0028823114202123096), ('virginica', 0.9971176508474883)]
Sample 4: [('setosa', 0.013209318664984725), ('versicolor', 0.7593991586796384), ('virginica', 0.22739152265537677)]
Sample 5: [('setosa', 0.0018885607572993624), ('versicolor', 0.7521357550798934), ('virginica', 0.24597568416280738)]
```

Practical No 3

Aim:- Implement Multinomial Logistic Regression (Iris Dataset)

Theory:

A **multinomial logistic regression** (or multinomial regression for short) is used when the outcome variable being predicted is nominal and has more than two categories that do not have a given rank or order. This model can be used with any number of independent variables that are categorical or continuous.

CODE:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

# Load the Iris dataset from sklearn
iris = datasets.load_iris()

# Convert the data into a DataFrame
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
iris_df['target'] = iris.target
```

iris_df.shape

```
(150, 5)
```

iris_df.columns

```
Index(['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)',  
      'petal width (cm)', 'target'],  
      dtype='object')
```

iris_df.info()

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 150 entries, 0 to 149  
Data columns (total 5 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   sepal length (cm)      150 non-null    float64  
1   sepal width (cm)       150 non-null    float64  
2   petal length (cm)      150 non-null    float64  
3   petal width (cm)       150 non-null    float64  
4   target                 150 non-null    int32  
dtypes: float64(4), int32(1)  
memory usage: 5.4 KB
```

```
iris_df.describe()
```

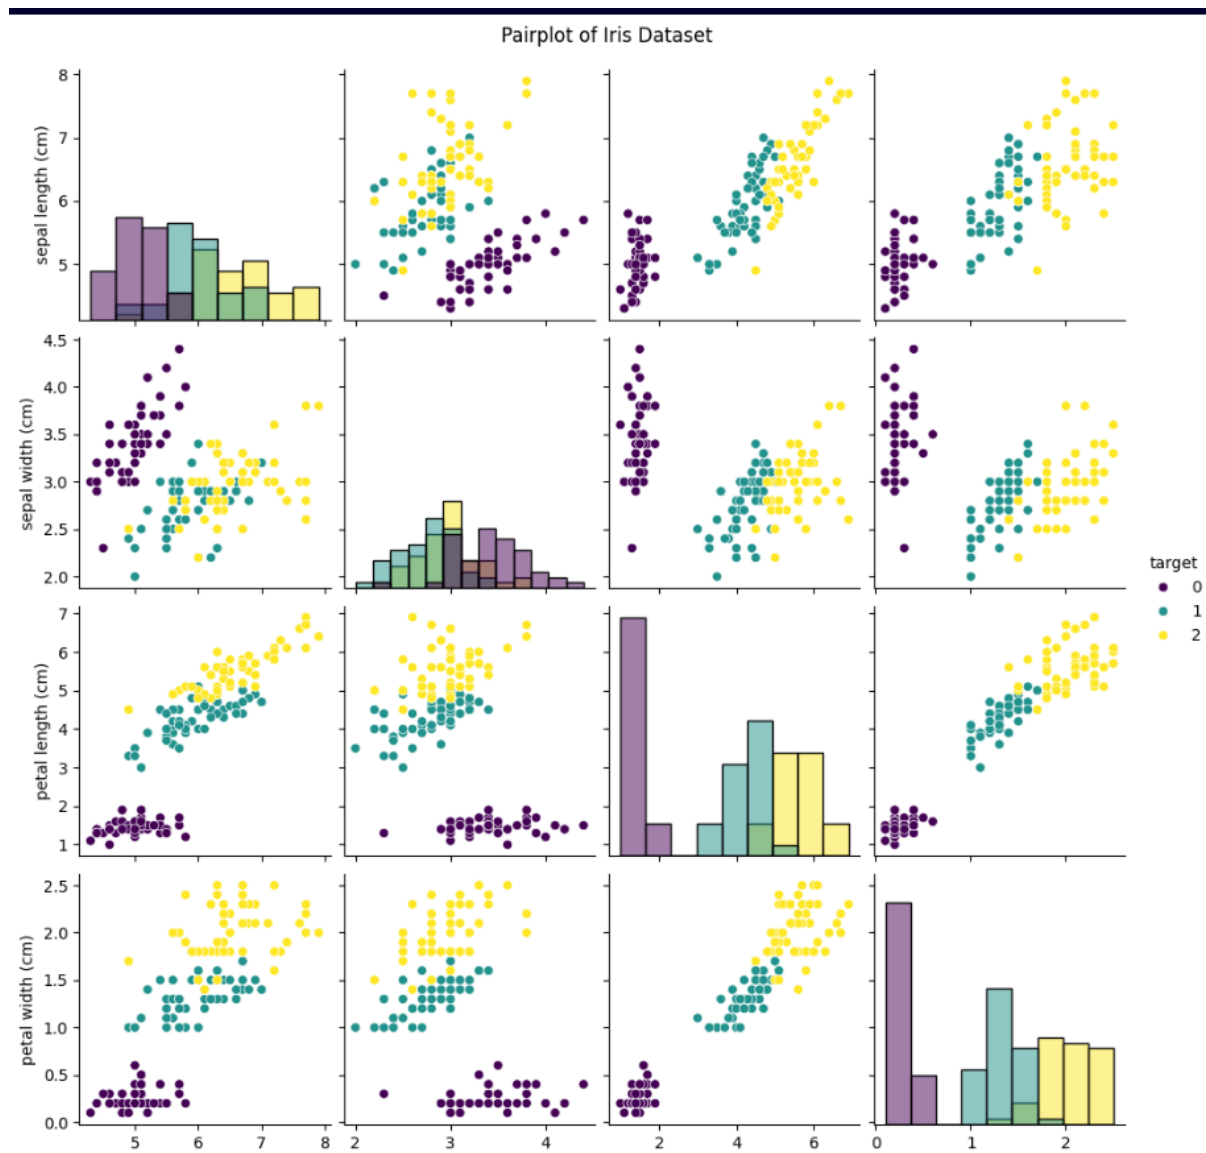
	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000
max	7.900000	4.400000	6.900000	2.500000	2.000000

EDA: Pairplot to visualize relationships between features

```
sns.pairplot(iris_df, hue='target', palette='viridis', diag_kind='hist')
```

```
plt.suptitle('Pairplot of Iris Dataset', y=1.02)
```

```
plt.show()
```



Other Plots

```
plt.figure(figsize=(12, 6))
```

Boxplot

```
plt.subplot(1, 2, 1)
```

```
sns.boxplot(x='target', y='sepal length (cm)', data=iris_df)
```

```
plt.title('Boxplot of Sepal Length by Target')
```

Violin Plot

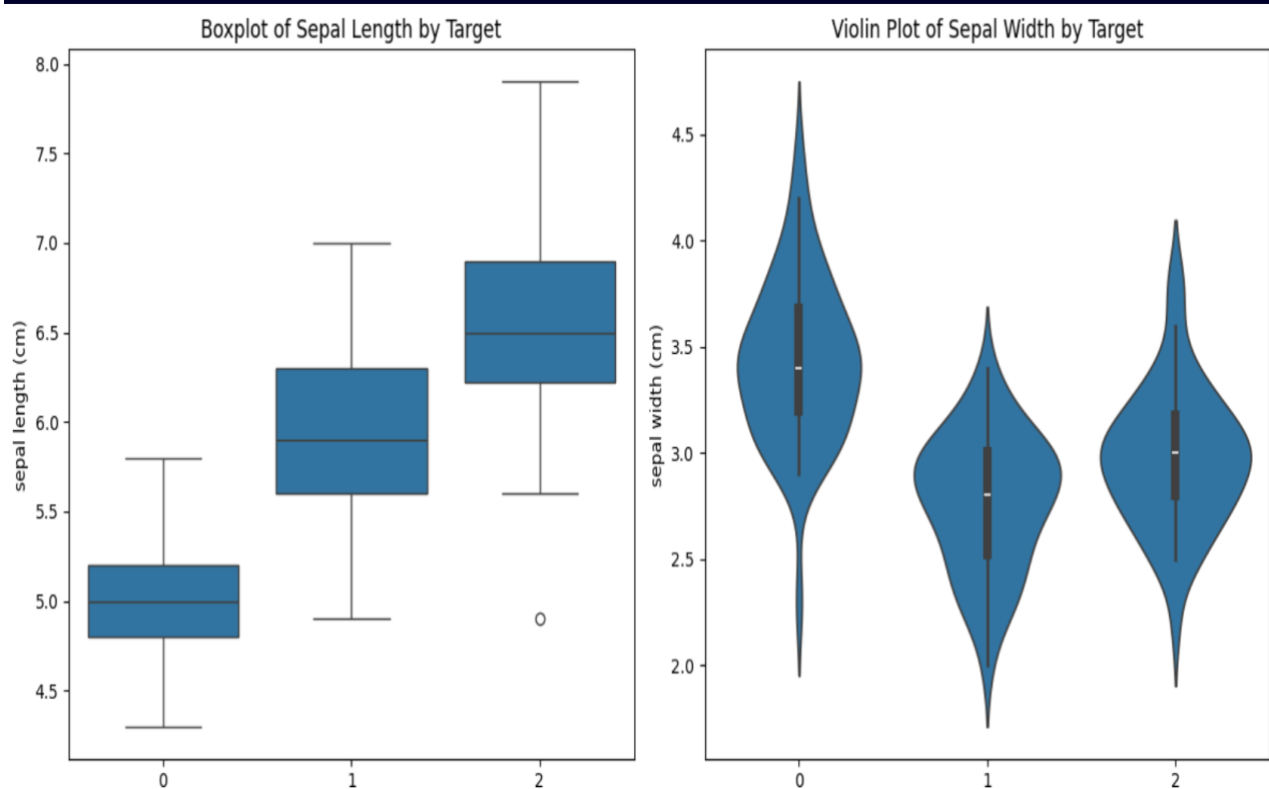
plt.subplot(1, 2, 2)

sns.violinplot(x='target', y='sepal width (cm)', data=iris_df)

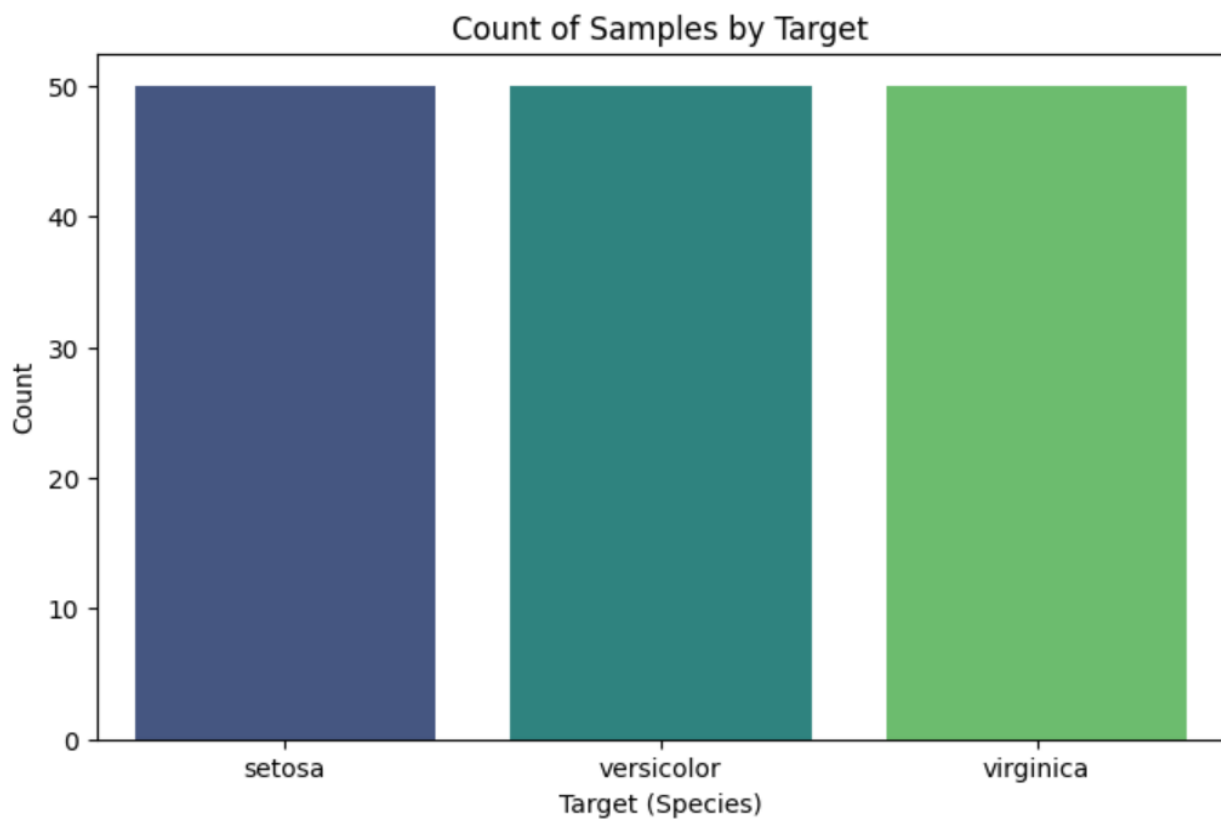
plt.title('Violin Plot of Sepal Width by Target')

plt.tight_layout()

plt.show()




```
# Count Plot (Bar Plot)
plt.figure(figsize=(8, 5))
sns.countplot(x='target', data=iris_df, palette='viridis')
plt.title('Count of Samples by Target')
plt.xlabel('Target (Species)')
plt.ylabel('Count')
plt.xticks(ticks=[0, 1, 2], labels=iris.target_names)
plt.show()
```



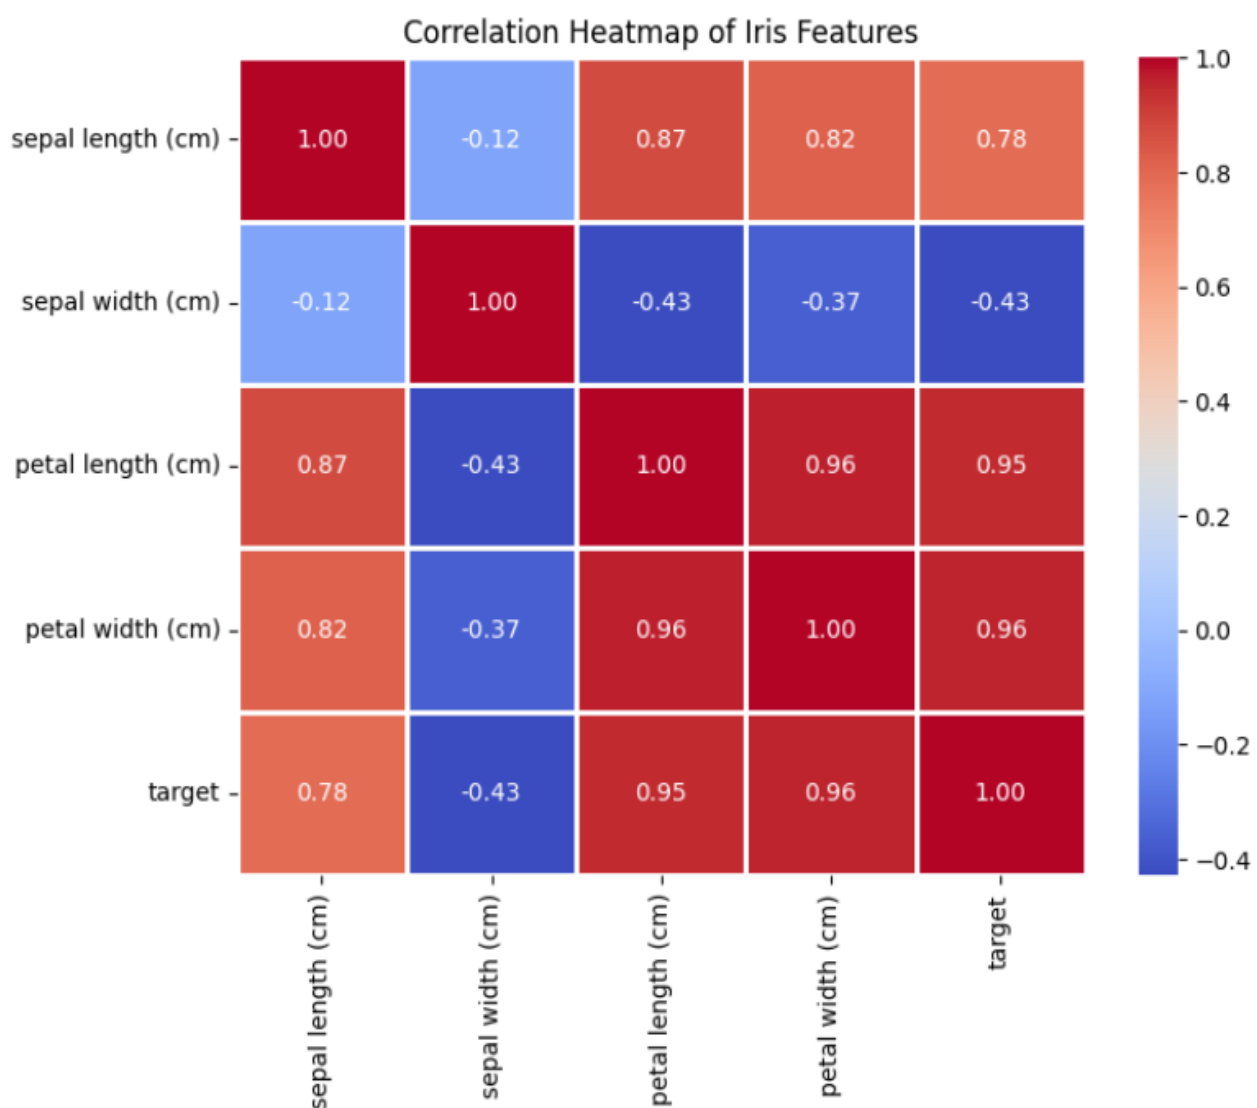
Correlation Heatmap

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(iris_df.corr(), annot=True, cmap='coolwarm', fmt='.2f',  
linewidths=1, linecolor='white')
```

```
plt.title('Correlation Heatmap of Iris Features')
```

```
plt.show()
```



Split the dataset into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

Standardize features by removing the mean and scaling to unit variance

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

Initialize the Logistic Regression model

```
model = LogisticRegression(multi_class="multinomial")
```

Train the model on the training data

```
model.fit(X_train, y_train)
```

Predict on the test data

```
y_pred = model.predict(X_test)
```

Calculate training and testing accuracy

```
train_accuracy = accuracy_score(y_train, model.predict(X_train))
```

```
test_accuracy = accuracy_score(y_test, y_pred)
```

```
print("Training Accuracy:", train_accuracy)
```

```
print("Testing Accuracy:", test_accuracy)
```

```
Training Accuracy: 0.9666666666666667
Testing Accuracy: 1.0
```

Create a confusion matrix

```
conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(conf_matrix)
```

```
Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

Print classification report

```
print("\nClassification Report:")
print(classification_report(y_test, y_pred,
target_names=iris.target_names))
```

```
Classification Report:
              precision    recall  f1-score   support

   setosa         1.00        1.00        1.00         10
  versicolor      1.00        1.00        1.00          9
   virginica      1.00        1.00        1.00         11
```

accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Print probabilities of classification for the first few samples in the test set

print("\nProbabilities of Classification:")

probabilities = model.predict_proba(X_test[:5])

for i, prob in enumerate(probabilities):

print(f"Sample {i+1}: {list(zip(iris.target_names, prob))}")

Probabilities of Classification:

Sample 1: [('setosa', 0.011457196118264058), ('versicolor', 0.8759785262009813), ('virginica', 0.11256427768075462)]
Sample 2: [('setosa', 0.9644113023352995), ('versicolor', 0.03558828637713127), ('virginica', 4.1128756931417627e-07)]
Sample 3: [('setosa', 3.7732299440075534e-08), ('versicolor', 0.0028823114202123196), ('virginica', 0.9971176508474883)]
Sample 4: [('setosa', 0.0132093186649847), ('versicolor', 0.7593991586796384), ('virginica', 0.22739152265537677)]
Sample 5: [('setosa', 0.001888560757299356), ('versicolor', 0.7521357550798935), ('virginica', 0.2459756841628072)]

Practical No 4

Aim: Implement SVM Classifier (Iris Datasets)

Theory:

A **support vector machine** (SVM) is defined as a machine learning algorithm that uses supervised learning models to solve complex classification, regression, and outlier detection problems by performing optimal data transformations that determine boundaries between data points based on predefined classes, labels.

CODE:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import classification_report, accuracy_score

# Load the Iris dataset from sklearn
iris = datasets.load_iris()

# Convert the data into a DataFrame
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
iris_df['target'] = iris.target
```

Exploratory Data Analysis (EDA)

print(iris_df.describe()) # Summary statistics

print(iris_df.head()) # View first few rows

	sepal length (cm)	sepal width (cm)	petal length (cm)	\
count	150.000000	150.000000	150.000000	
mean	5.843333	3.057333	3.758000	
std	0.828066	0.435866	1.765298	
min	4.300000	2.000000	1.000000	
25%	5.100000	2.800000	1.600000	
50%	5.800000	3.000000	4.350000	
75%	6.400000	3.300000	5.100000	
max	7.900000	4.400000	6.900000	

	petal width (cm)	target
count	150.000000	150.000000
mean	1.199333	1.000000
std	0.762238	0.819232
min	0.100000	0.000000
25%	0.300000	0.000000
50%	1.300000	1.000000
75%	1.800000	2.000000
max	2.500000	2.000000

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

Visualization (pairplot for all features)

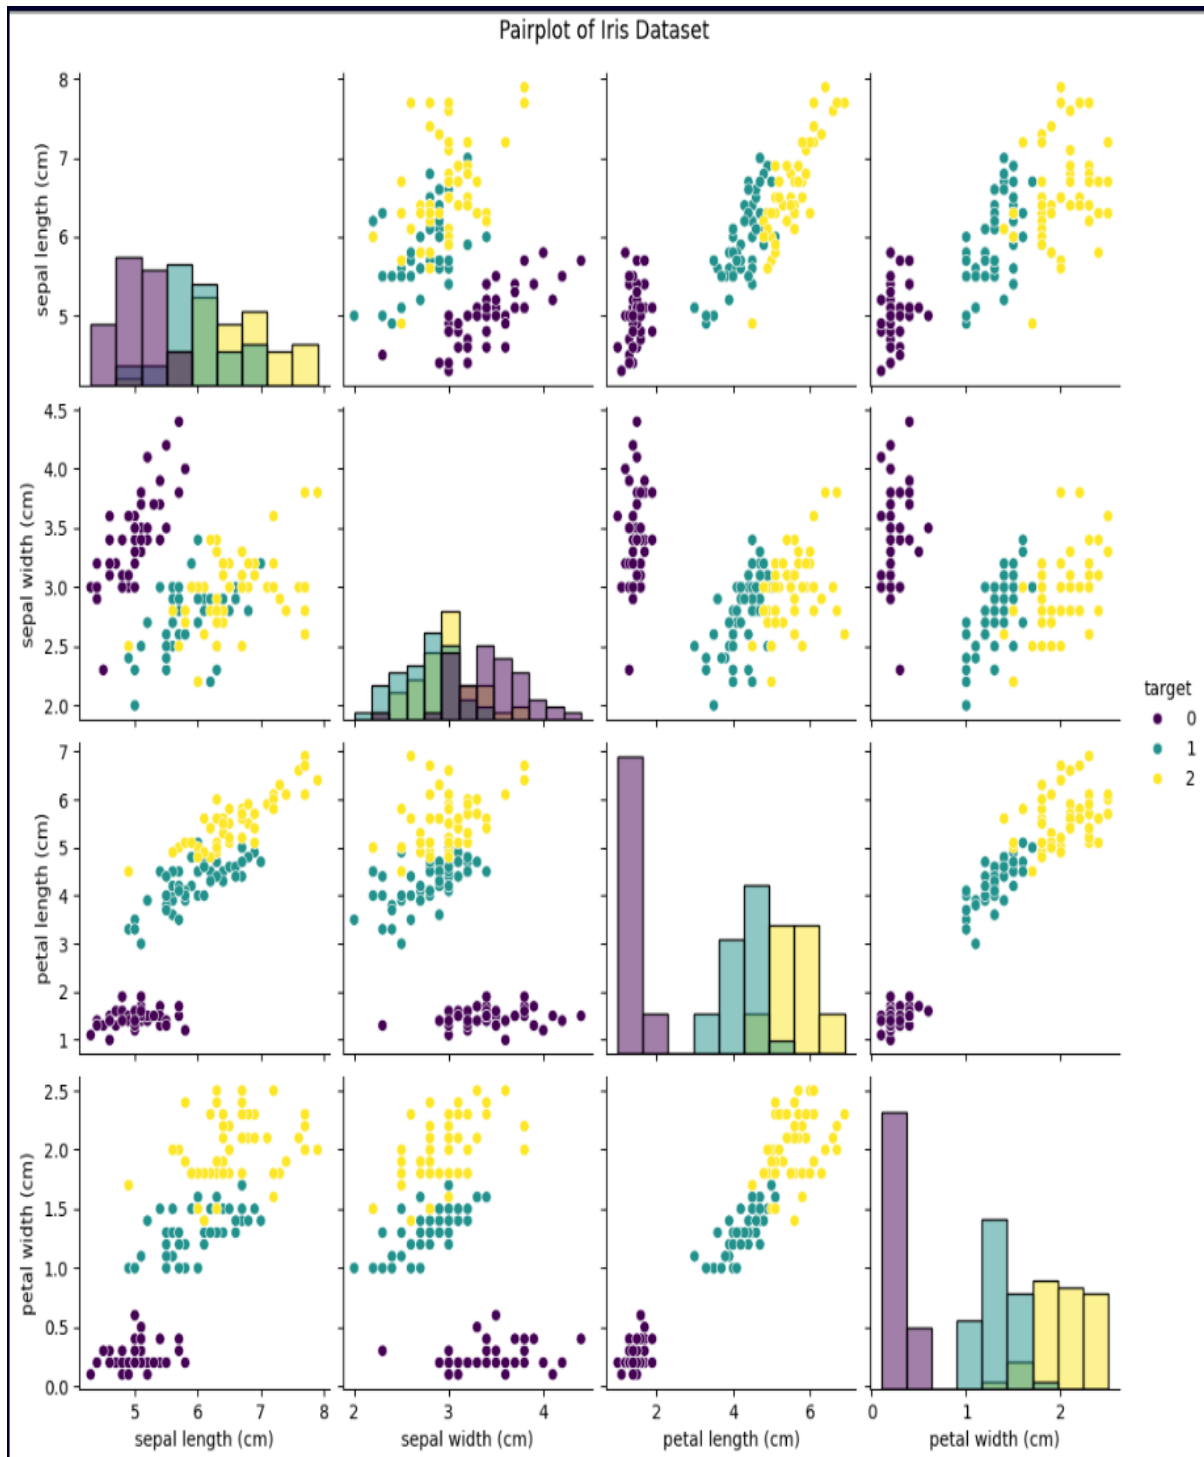
import seaborn as sns

EDA: Pairplot to visualize relationships between features

sns.pairplot(iris_df, hue='target', palette='viridis', diag_kind='hist')

plt.suptitle("Pairplot of Iris Dataset", y=1.02)

`plt.show()`



Separate features (X) and target (y) from the DataFrame

X = iris_df.drop('target', axis=1) # Features

y = iris_df['target'] # Target (labels)

Split the dataset into training and testing sets

**X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)**

Create the SVM model

clf = SVC(kernel='linear') # Experiment with different kernels (e.g., 'rbf')

Train the model

clf.fit(X_train, y_train)

Make predictions on the testing set

y_pred = clf.predict(X_test)

Evaluate model performance

print(classification_report(y_test, y_pred))

print("Training Accuracy:", accuracy_score(y_train, clf.predict(X_train)))

print("Testing Accuracy:", accuracy_score(y_test, y_pred))

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30
Training Accuracy: 0.975				
Testing Accuracy: 1.0				

HYPER PARAMETER TUNING

```

from sklearn.model_selection import GridSearchCV

# Define a parameter grid to explore
param_grid = {'kernel': ['linear', 'rbf'],
              'C': [0.01, 0.1, 1, 10, 100]}

# Create the GridSearchCV object
grid_search = GridSearchCV(SVC(), param_grid, cv=5) # 5-fold cross-validation

# Fit the grid search to the training data
grid_search.fit(X_train, y_train)

# Get the best model and its parameters
best_model = grid_search.best_estimator_
best_params = grid_search.best_params_
print(best_params)

# Use the best model for prediction and evaluation
y_pred = best_model.predict(X_test)

```

```
print(classification_report(y_test, y_pred))
print("Testing Accuracy:", accuracy_score(y_test, y_pred))
```

```
{'C': 1, 'kernel': 'linear'}
      precision    recall  f1-score   support

     0       1.00      1.00      1.00        10
     1       1.00      1.00      1.00         9
     2       1.00      1.00      1.00        11

 accuracy          1.00          1.00          1.00        30
 macro avg          1.00          1.00          1.00        30
weighted avg          1.00          1.00          1.00        30

Testing Accuracy: 1.0
```

Practical No 5

Aim: Train and fine-tune a Decision Tree for Moon dataset.

Theory:

A **decision tree** is a decision support hierarchical model that uses a tree-like model of decisions and their possible consequences, including chance event outcomes, resource costs, and utility.

CODE:

```
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.tree import plot_tree

# Generate moons data
X, y = make_moons(n_samples=1000, noise=0.3)

# 1. Data Shape and Description
print("Data Shape:", X.shape)
print("Description of first 5 rows:")
print(X[:5])
print("Description of target variable:")
print(pd.Series(y).value_counts()) # Convert y to pandas Series
```

```
Data Shape: (1000, 2)
Description of first 5 rows:
[[ 0.0855163  1.30521226]
 [-0.05544742  0.2908183 ]
 [ 1.60334944 -0.61151162]
 [ 1.38402564 -0.60404707]
 [ 1.51886676  0.02146531]]
Description of target variable:
0    500
1    500
dtype: int64
```


2. Check for Missing Values

```
import numpy as np
print("Missing values in features:", np.isnan(X).sum(axis=0))
```

```
Missing values in features: [0 0]
```

3. Visualize the moons data

```
plt.scatter(X[:, 0], X[:, 1], c=y)
plt.title("Moons Dataset")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

```
# Define the decision tree classifier
```

```
clf = DecisionTreeClassifier()
```

5. Hyperparameter Tuning with GridSearchCV:

We'll tune two important hyperparameters for decision trees:

`max_depth`: Maximum depth of the tree. `min_samples_split`: Minimum number of samples required to split a node.

```
# Define hyperparameter grid
```

```
param_grid = {  
    'max_depth': [2, 3, 4, 5],  
    'min_samples_split': [2, 5, 10]
```

```
}
```

```
# Create GridSearchCV object
```

```
grid_clf = GridSearchCV(clf, param_grid, scoring='accuracy')
```

```
# Train the model
```

```
grid_clf.fit(X_train, y_train)
```

```
# Get the best model
```

```
best_model = grid_clf.best_estimator_
```

```
# Print the best hyperparameters
```

```
print("Best Hyperparameters:", grid_clf.best_params_)
```

```
Best Hyperparameters: {'max_depth': 5, 'min_samples_split': 2}
```

```
# Predict on test set
```

```
y_pred = best_model.predict(X_test)
```

```
# Calculate accuracy
```

```
from sklearn.metrics import accuracy_score
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print("Test Accuracy:", accuracy)
```


```
Test Accuracy: 0.9
```

```
#
```

Visualize the decision tree

plot_tree(best_model)

plt.show()



Practical No 6

Aim: Train an SVM regression on the California Housing Dataset

Theory:

Support Vector Regression is an extension of SVM which introduces a region, named tube, around the function to optimize with the aim of finding the tube that best approximates the continuous-valued function, while minimizing the prediction error, that is, the difference between the predicted and the true class label.

CODE:

```
import pandas as pd

from sklearn.datasets import fetch_california_housing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score

import seaborn as sns # For visualization
import matplotlib.pyplot as plt # For visualization

# Load data
data = fetch_california_housing()
X = pd.DataFrame(data.data, columns=data.feature_names)
y = data.target
```

X.columns

```
Index(['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup',
      'Latitude', 'Longitude'],
      dtype='object')
```

y

```
array([4.526, 3.585, 3.521, ..., 0.923, 0.847, 0.894])
```

Print basic info about the data

`print(X.describe())` # Summary statistics

`print(X.head())` # View first few rows

```
count      MedInc      HouseAge      AveRooms      AveBedrms      Population  \
mean         3.870671      28.639486         5.429000         1.096675      1425.476744
std          1.899822      12.585558         2.474173         0.473911      1132.462122
min           0.499900         1.000000         0.846154         0.333333         3.000000
25%          2.563400      18.000000         4.440716         1.006079         787.000000
50%          3.534800      29.000000         5.229129         1.048780      1166.000000
75%          4.743250      37.000000         6.052381         1.099526      1725.000000
max          15.000100      52.000000      141.909091         34.066667     35682.000000

count      AveOccup      Latitude      Longitude
mean         3.070655      35.631861      -119.569704
std          10.386050         2.135952         2.003532
min           0.692308      32.540000      -124.350000
25%          2.429741      33.930000      -121.800000
50%          2.818116      34.260000      -118.490000
75%          3.282261      37.710000      -118.010000
max          1243.333333      41.950000      -114.310000

   MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  \
0   8.3252     41.0   6.984127   1.023810         322.0   2.555556     37.88
1   8.3014     21.0   6.238137   0.971880         2401.0   2.109842     37.86
2   7.2574     52.0   8.288136   1.073446          496.0   2.802260     37.85
3   5.6431     52.0   5.817352   1.073059          558.0   2.547945     37.85
4   3.8462     52.0   6.281853   1.081081          565.0   2.181467     37.85
...
1  -122.22
2  -122.24
3  -122.25
4  -122.25
```

Check for missing values

print("Missing values:", X.isnull().sum())

```
Missing values: MedInc      0
HouseAge      0
AveRooms      0
AveBedrms     0
Population    0
AveOccup      0
Latitude      0
Longitude     0
dtype: int64
```

Exploratory Data Analysis (EDA)

Visualize the distribution of the target variable (median house value)


sns.distplot(y)

plt.xlabel("Median House Value")

plt.ylabel("Density")

plt.title("Distribution of Median House Values")

plt.show()



Data Preprocessing

Scale features (SVM regressor is sensitive to feature scales)

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)

Split Data into Training and Testing Sets

**X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)**

Train the SVM Regressor

svr = SVR(kernel='rbf') # Experiment with 'linear' or other kernels

svr.fit(X_train, y_train)

Make Predictions and Evaluate Performance

```
y_pred = svr.predict(X_test)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print("Mean Squared Error:", mse)
```

```
print("R-squared:", r2)
```

```
Mean Squared Error: 0.3551984619989417  
R-squared: 0.7289407597956463
```

```
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.svm import SVR
```

Define a parameter grid to explore

```
param_grid = {
```

```
    'kernel': ['linear', 'rbf'], # Experiment with different kernels
```

```
    'C': [0.01, 0.1, 1, 10, 100], # Regularization parameter
```

```
    'gamma': [0.001, 0.01, 0.1, 1], # Gamma for RBF kernel (optional)
```

```
}
```

Create the GridSearchCV object

```
grid_search = GridSearchCV(SVR(), param_grid, cv=5) # 5-fold cross-validation
```

```
# Fit the grid search to the training data
```

```
grid_search.fit(X_train, y_train)
```

```
# Get the best model and its parameters
```

```
best_model = grid_search.best_estimator_
```

```
best_params = grid_search.best_params_
```

```
# Use the best model for prediction and evaluation
```

```
y_pred = best_model.predict(X_test)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

Practical No 7

Aim: Implement Batch Gradient Descent with Early Stopping for Softmax Regression.

Theory:

Batch gradient descent, also called vanilla gradient descent, calculates the error for each example within the training dataset, but only after all training examples have been evaluated does the model get updated. This whole process is like a cycle and it's called a training epoch.

In machine learning, **early stopping** is a form of regularization used to avoid overfitting when training a learner with an iterative method, such as gradient descent. Such methods update the learner so as to make it better fit the training data with each iteration.

SoftMax is particularly suited for multi-class classification problems, as it provides a clear and normalized probability distribution across all possible classes.

CODE:

```
import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

# Load the Iris dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target

# Standardize the features
scaler = StandardScaler()
```

```

X = scaler.fit_transform(X)

# Add a bias term (column of ones) to the data
X = np.c_[np.ones(X.shape[0]), X]

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

def softmax(logits):
    exp_logits = np.exp(logits - np.max(logits, axis=1, keepdims=True))
    return exp_logits / np.sum(exp_logits, axis=1, keepdims=True)

def compute_loss_and_gradients(X, y, theta):
    logits = X.dot(theta)
    y_proba = softmax(logits)
    m = X.shape[0]
    entropy_loss = -np.mean(np.log(y_proba[np.arange(m), y]))
    gradients = (1/m) * X.T.dot(y_proba - np.eye(np.max(y) + 1)[y])
    return entropy_loss, gradients

def predict(X, theta):
    logits = X.dot(theta)
    return np.argmax(softmax(logits), axis=1)

def softmax_regression(X_train, y_train, X_val, y_val, learning_rate=0.01,
n_epochs=1000, tol=1e-4, patience=5):
    n_inputs = X_train.shape[1]

```



```

n_outputs = np.max(y_train) + 1
theta = np.random.randn(n_inputs, n_outputs)

best_loss = np.inf
epochs_without_improvement = 0

for epoch in range(n_epochs):
    loss, gradients = compute_loss_and_gradients(X_train, y_train, theta)
    theta = theta - learning_rate * gradients

    val_loss, _ = compute_loss_and_gradients(X_val, y_val, theta)

    if val_loss < best_loss - tol:
        best_loss = val_loss
        epochs_without_improvement = 0
    else:
        epochs_without_improvement += 1

    if epochs_without_improvement >= patience:
        print(f"Early stopping at epoch {epoch}")
        break

return theta

# Split the training data into training and validation sets
X_train_split, X_val_split, y_train_split, y_val_split =
train_test_split(X_train, y_train, test_size=0.2, random_state=42)

```

Train the model

**theta = softmax_regression(X_train_split, y_train_split, X_val_split,
y_val_split)**

Test accuracy: 93.33%

Practical No 8

Aim: Implement MLP for Classification of Handwritten digits(MNIST datasets)

Theory:

The **MNIST database** (Modified National Institute of Standards and Technology database) is a large collection of handwritten digits. It has a training set of 60,000 examples, and a test set of 10,000 examples.

CODE:

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
import matplotlib.pyplot as plt

# Load and preprocess the MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()
print(x_train.shape)
x_train = x_train / 255.0
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)

# Define the CNN model architecture
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
```

```
MaxPooling2D((2, 2)),
Conv2D(64, (3, 3), activation='relu'),
MaxPooling2D((2, 2)),
Flatten(),
Dense(128, activation='relu'),
Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model using the training data
model.fit(x_train, y_train, epochs=5)

# Choose a single image from the test set
index = 0 # Replace with the index of the image you want to use
single_image = x_test[index]
input_image = np.expand_dims(single_image, axis=0)

# Get the predicted probabilities for the single image
predicted_probabilities = model.predict(input_image)

# Display the input image
plt.imshow(single_image, cmap='gray')
plt.title('Input Image')
```

```
plt.axis('off')
```

```
plt.show()
```

```
# Display the predicted probabilities
```

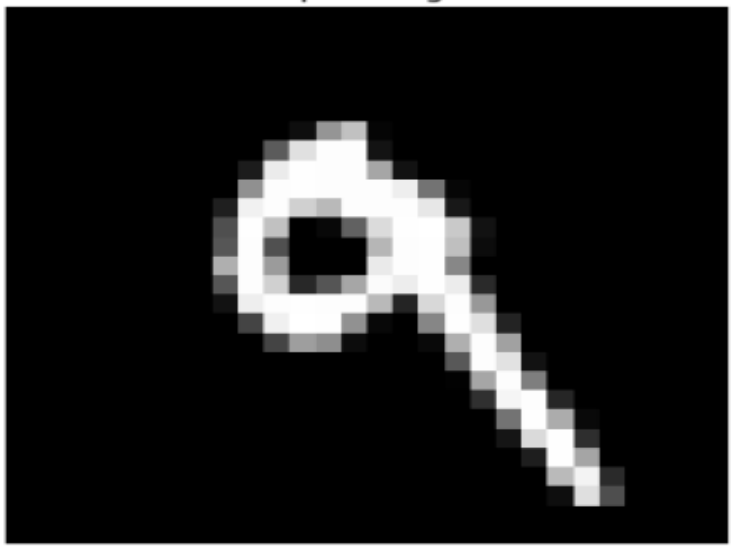
```
print("Predicted Probabilities:", predicted_probabilities)
```

```
# Get the predicted class (index with highest probability)
```

```
predicted_class = np.argmax(predicted_probabilities)
```

```
print("Predicted Class:", predicted_class)
```

```
(60000, 28, 28)
Epoch 1/5
1875/1875 ————— 30s 13ms/step - accuracy: 0.9015 - loss: 0.3057
Epoch 2/5
1875/1875 ————— 25s 13ms/step - accuracy: 0.9855 - loss: 0.0459
Epoch 3/5
1875/1875 ————— 25s 14ms/step - accuracy: 0.9911 - loss: 0.0286
Epoch 4/5
1875/1875 ————— 26s 14ms/step - accuracy: 0.9941 - loss: 0.0201
Epoch 5/5
1875/1875 ————— 25s 13ms/step - accuracy: 0.9954 - loss: 0.0145
1/1 ————— 0s 98ms/step
```



```
Predicted Probabilities: [[0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]]
Predicted Class: 9
```

Practical No 9

Aim: Classification of Image of clothing using Tensorflow(Fashion MNIST dataset)

Theory:

TensorFlow is a free and open-source software library for machine learning and artificial intelligence. It can be used across a range of tasks but has a particular focus on training and inference of deep neural networks. It was developed by the Google Brain team for Google's internal use in research and production

CODE:

```
import tensorflow as tf

from tensorflow.keras import datasets, layers, models

import matplotlib.pyplot as plt

# Load the Fashion MNIST dataset

(train_images, train_labels), (test_images, test_labels) =
datasets.fashion_mnist.load_data()

# Normalize the images to a range of 0 to 1

train_images, test_images = train_images / 255.0, test_images / 255.0

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
```

```

layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation='relu'),
layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Reshape the data to include the channel dimension
train_images = train_images.reshape((train_images.shape[0], 28, 28, 1))
test_images = test_images.reshape((test_images.shape[0], 28, 28, 1))

# Train the model
history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))

```

```

Epoch 1/10
1875/1875 — 27s 13ms/step - accuracy: 0.7506 - loss: 0.6901 - val_accuracy: 0.8742 - val_loss: 0.3552
Epoch 2/10
1875/1875 — 40s 12ms/step - accuracy: 0.8773 - loss: 0.3322 - val_accuracy: 0.8817 - val_loss: 0.3256
Epoch 3/10
1875/1875 — 41s 13ms/step - accuracy: 0.8981 - loss: 0.2773 - val_accuracy: 0.8768 - val_loss: 0.3421
Epoch 4/10
1875/1875 — 24s 13ms/step - accuracy: 0.9087 - loss: 0.2479 - val_accuracy: 0.9012 - val_loss: 0.2738
Epoch 5/10
1875/1875 — 24s 13ms/step - accuracy: 0.9193 - loss: 0.2200 - val_accuracy: 0.8977 - val_loss: 0.2765
Epoch 6/10
1875/1875 — 42s 13ms/step - accuracy: 0.9253 - loss: 0.1996 - val_accuracy: 0.8988 - val_loss: 0.2741
Epoch 7/10
1875/1875 — 24s 13ms/step - accuracy: 0.9322 - loss: 0.1795 - val_accuracy: 0.9077 - val_loss: 0.2666
Epoch 8/10
1875/1875 — 25s 13ms/step - accuracy: 0.9386 - loss: 0.1655 - val_accuracy: 0.9030 - val_loss: 0.2866
Epoch 9/10
1875/1875 — 41s 13ms/step - accuracy: 0.9449 - loss: 0.1507 - val_accuracy: 0.9107 - val_loss: 0.2738
Epoch 10/10
1875/1875 — 27s 15ms/step - accuracy: 0.9476 - loss: 0.1382 - val_accuracy: 0.9097 - val_loss: 0.2876

```

```
plt.figure(figsize=(12, 4))
```

```
plt.subplot(1, 2, 1)
```

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
```

```
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Accuracy')
```

```
plt.legend()
```

```
plt.title('Accuracy')
```

```
plt.subplot(1, 2, 2)
```

```
plt.plot(history.history['loss'], label='Training Loss')
```

```
plt.plot(history.history['val_loss'], label='Validation Loss')
```


```
plt.xlabel('Epoch')
```

```
plt.ylabel('Loss')
```

```
plt.legend()
```

```
plt.title('Loss')
```

```
plt.show()
```




```

import numpy as np
predictions = model.predict(test_images)

# Define a function to plot the images and predictions
def plot_image(predictions_array, true_label, img):
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])

    plt.imshow(img, cmap=plt.cm.binary)

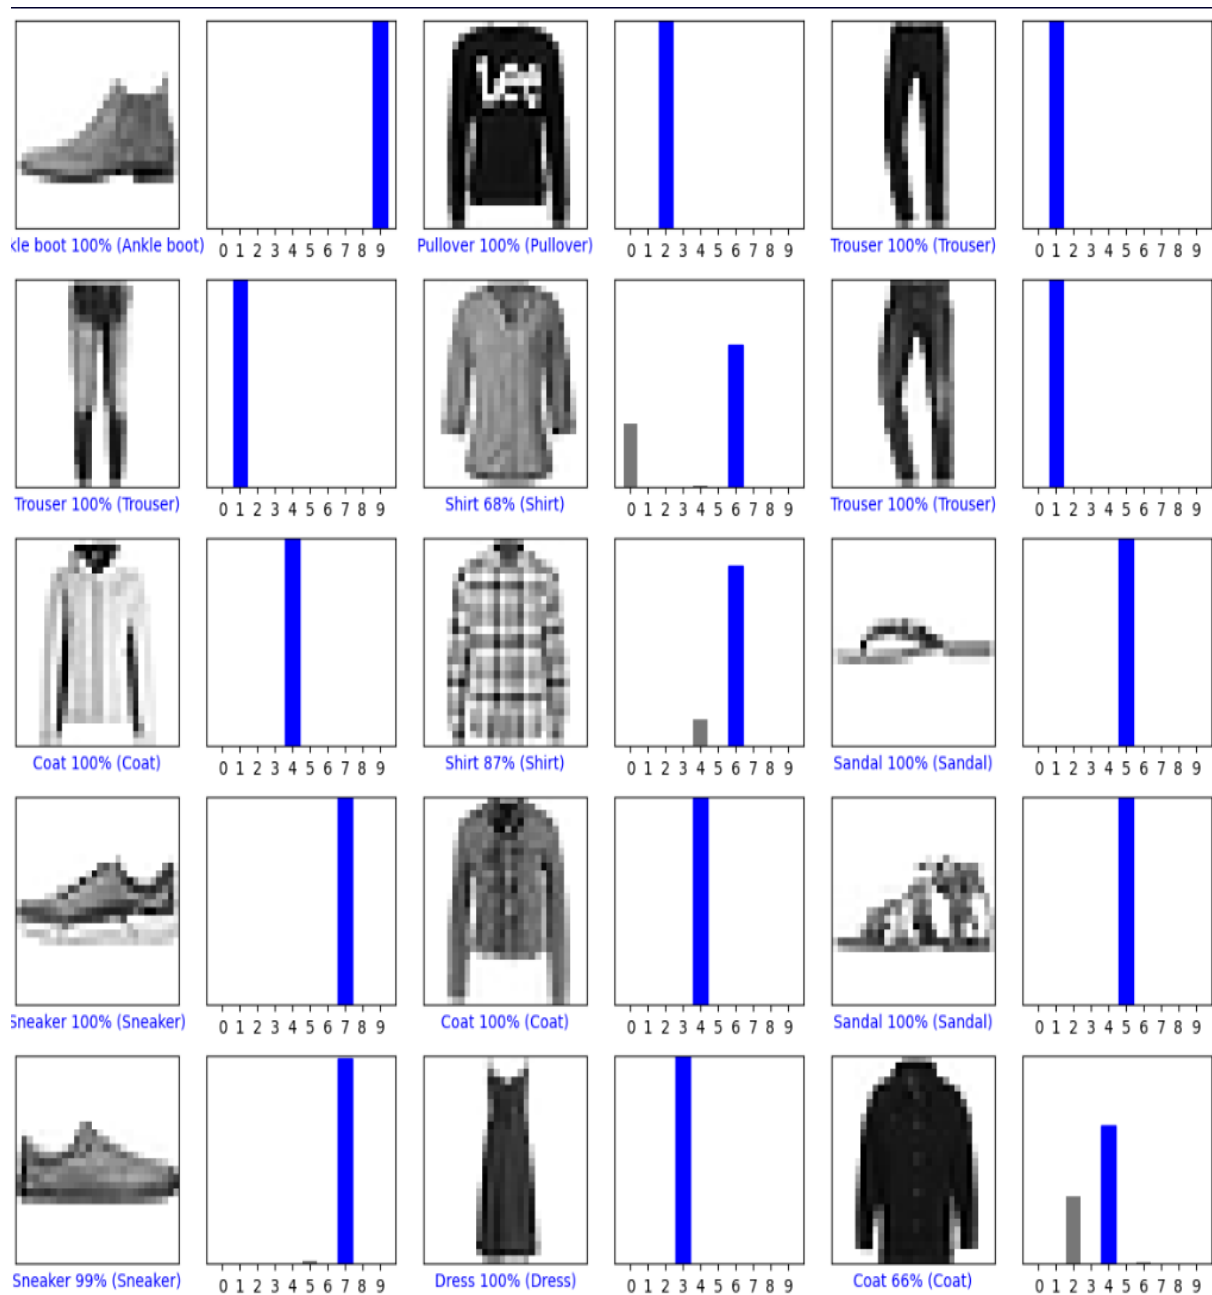
    predicted_label = np.argmax(predictions_array)
    if predicted_label == true_label:
        color = 'blue'
    else:
        color = 'red'

    plt.xlabel(f'{class_names[predicted_label]}
{100*np.max(predictions_array):2.0f}% ({class_names[true_label]})',
color=color)

def plot_value_array(predictions_array, true_label):
    plt.grid(False)
    plt.xticks(range(10))
    plt.yticks([])
    thisplot = plt.bar(range(10), predictions_array, color="#777777")
    plt.ylim([0, 1])

```

```
predicted_label = np.argmax(predictions_array)  
  
thisplot[predicted_label].set_color('red')  
thisplot[true_label].set_color('blue')  
  
# Define the class names  
class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',  
                'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']  
  
# Plot the first 15 test images, their predicted labels, and the true labels  
# Color correct predictions in blue and incorrect predictions in red  
num_rows = 5  
num_cols = 3  
num_images = num_rows * num_cols  
plt.figure(figsize=(2*2*num_cols, 2*num_rows))  
for i in range(num_images):  
    plt.subplot(num_rows, 2*num_cols, 2*i+1)  
    plot_image(predictions[i], test_labels[i], test_images[i].reshape(28, 28))  
    plt.subplot(num_rows, 2*num_cols, 2*i+2)  
    plot_value_array(predictions[i], test_labels[i])  
plt.tight_layout()  
plt.show()
```



Practical No 10

Aim: Implement Regression to predict fuel efficiency using TensorFlow (Auto MPG dataset).

Theory:

Regression is a statistical method used in finance, investing, and other disciplines that attempts to determine the strength and character of the relationship between a dependent variable and one or more independent variables. Linear regression is the most common form of this technique.

TensorFlow is a free and open-source software library for machine learning and artificial intelligence. It can be used across a range of tasks but has a particular focus on training and inference of deep neural networks. It was developed by the Google Brain team for Google's internal use in research and production

CODE:

```
import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Load the dataset
url = "auto-mpg.csv" # Replace with your CSV file path
column_names = ['MPG', 'Cylinders', 'Displacement', 'Horsepower',
'Weight', 'Acceleration', 'Model Year', 'Origin']
```

```
dataset = pd.read_csv(url, names=column_names, na_values='?',
comment='\t', sep=',', skipinitialspace=True)
```

`dataset.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 399 entries, mpg to 31
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MPG                   399 non-null    object
1   Cylinders             399 non-null    object
2   Displacement          393 non-null    object
3   Horsepower            399 non-null    object
4   Weight                399 non-null    object
5   Acceleration          399 non-null    object
6   Model Year            399 non-null    object
7   Origin                399 non-null    object
dtypes: object(8)
memory usage: 28.1+ KB
```

`dataset.head()`

	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model Year	Origin
mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
18	8	307	130	3504	12	70	1	chevrolet chevelle malibu
15	8	350	165	3693	11.5	70	1	buick skylark 320
18	8	318	150	3436	11	70	1	plymouth satellite
16	8	304	150	3433	12	70	1	amc rebel sst

`dataset.describe()`

dataset.describe()								
	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model Year	Origin
count	399	399	393	399	399	399	399	399
unique	6	83	94	352	96	14	4	306
top	4	97	150	1985	14.5	73	1	ford pinto
freq	204	21	22	4	23	40	249	6

Drop rows with missing values

dataset = dataset.dropna()

Convert columns to appropriate numeric data types

**for column in ['MPG', 'Cylinders', 'Displacement', 'Horsepower', 'Weight',
'Acceleration', 'Model Year']:**

dataset[column] = pd.to_numeric(dataset[column], errors='coerce')

Check for NaN values in the dataset

print("NaN values before dropping: \n", dataset.isnull().sum())

Drop any rows with NaN values

dataset = dataset.dropna()

Check again for NaN values to confirm

print("NaN values after dropping: \n", dataset.isnull().sum())

Convert 'Origin' to string for one-hot encoding

dataset['Origin'] = dataset['Origin'].astype(str)

Convert categorical 'Origin' column to one-hot encoding

**dataset = pd.get_dummies(dataset, columns=['Origin'], prefix='',
prefix_sep='')**

```
NaN values before dropping:
```

```
MPG          1
Cylinders    1
Displacement 1
Horsepower   1
Weight       1
Acceleration 1
Model Year   1
Origin       0
```

```
dtype: int64
```

```
NaN values after dropping:
```

```
MPG          0
Cylinders    0
Displacement 0
Horsepower   0
Weight       0
Acceleration 0
Model Year   0
Origin       0
```

```
dtype: int64
```

Separate features and labels

```
train_features = train_dataset.copy()
```

```
test_features = test_dataset.copy()
```

```
train_labels = train_features.pop('MPG')
```

```
test_labels = test_features.pop('MPG')
```

Check for NaN values in the dataset

```
assert not train_features.isnull().any().any(), "There are NaN values in the training features"
```

```
assert not test_features.isnull().any().any(), "There are NaN values in the test features"
```

```
assert not train_labels.isnull().any(), "There are NaN values in the training labels"
```

```
assert not test_labels.isnull().any(), "There are NaN values in the test labels"
```

```
# Normalize the features
```

```
scaler = StandardScaler()
```

```
train_features = scaler.fit_transform(train_features)
```

```
test_features = scaler.transform(test_features)
```

```
def build_model():
```

```
    model = models.Sequential([  
        layers.Dense(64, activation='relu',  
input_shape=[train_features.shape[1]],  
        layers.Dense(64, activation='relu'),  
        layers.Dense(1)  
])  
    return model
```

```
model = build_model()
```

```
model.compile(optimizer='adam',
```

```
    loss='mse',
```

```
    metrics=['mae', 'mse'])
```

```
history = model.fit(train_features, train_labels,  
    epochs=100, validation_split=0.2, verbose=0)
```

```
# Plot training history
```

```
hist = pd.DataFrame(history.history)
```

```
hist['epoch'] = history.epoch
```



```
plt.figure(figsize=(12, 4))
```

```
plt.subplot(1, 2, 1)
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Mean Abs Error [MPG]')
```

```
plt.plot(hist['epoch'], hist['mae'], label='Train Error')
```

```
plt.plot(hist['epoch'], hist['val_mae'], label='Val Error')
```

```
plt.legend()
```

```
plt.subplot(1, 2, 2)
```

```
plt.xlabel('Epoch')
```


```
plt.ylabel('Mean Square Error [MPG^2]')
```

```
plt.plot(hist['epoch'], hist['mse'], label='Train Error')
```

```
plt.plot(hist['epoch'], hist['val_mse'], label='Val Error')
```

```
plt.legend()
```

```
plt.show()
```



```
test_loss, test_mae, test_mse = model.evaluate(test_features, test_labels,  
verbose=2)
```

```
print(f'\nTest MAE: {test_mae:.2f} MPG')
```

```
3/3 - 0s - loss: 5.7426 - mae: 1.9889 - mse: 5.7426 - 52ms/epoch - 17ms/step
```

```
Test MAE: 1.99 MPG
```

```
test_predictions = model.predict(test_features).flatten()
```

```
plt.scatter(test_labels, test_predictions)
```

```
plt.xlabel('True Values [MPG]')
```

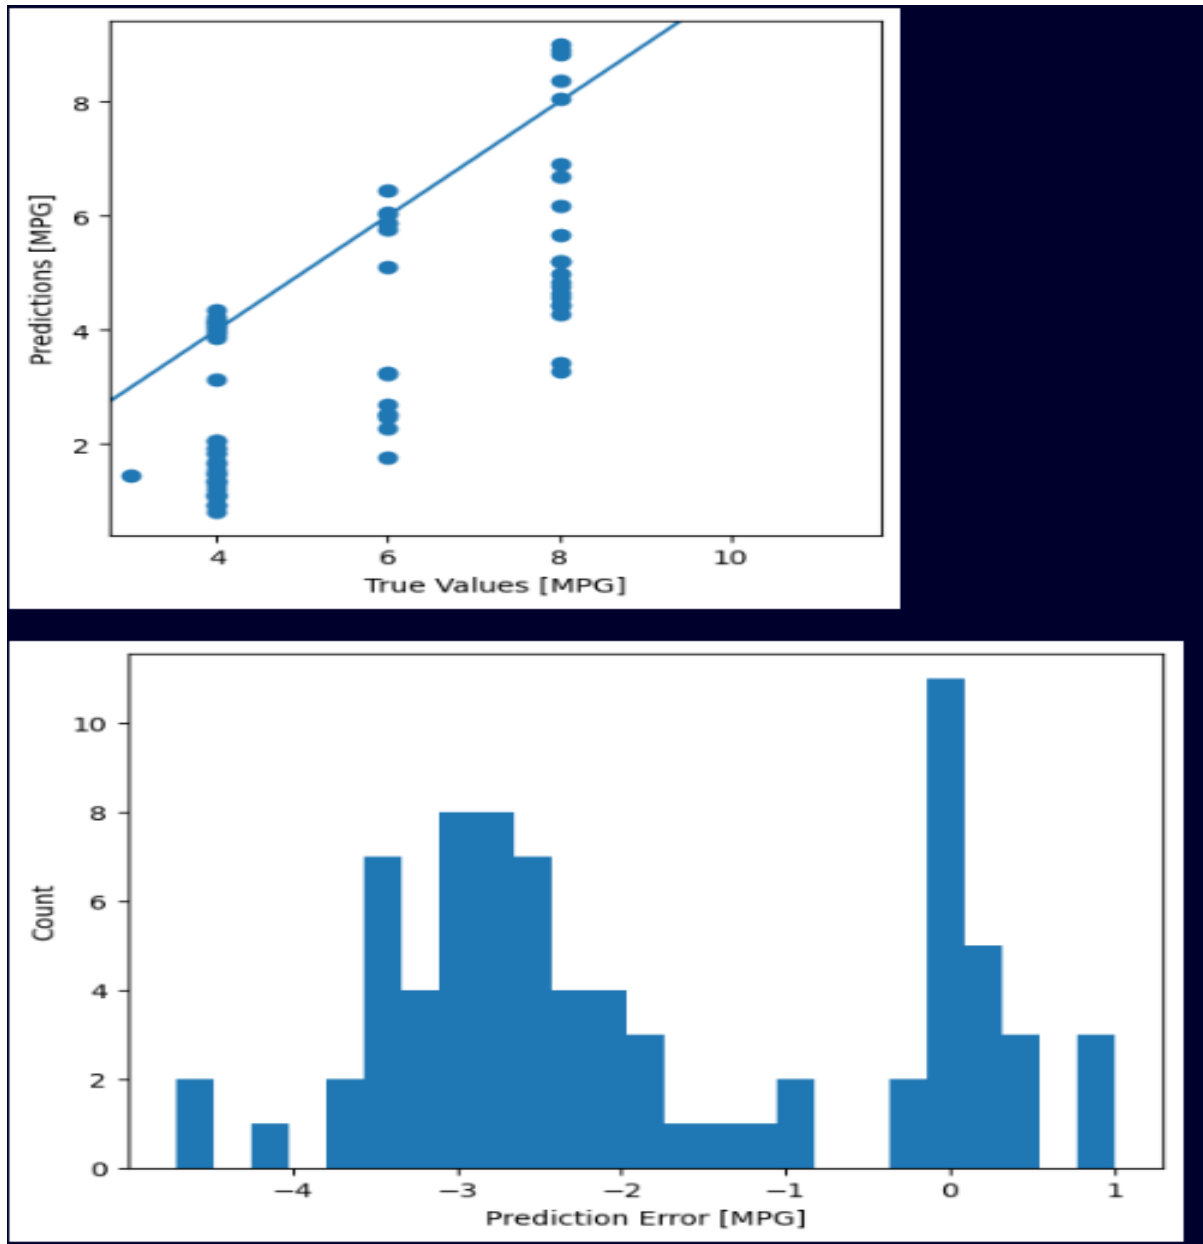
```
plt.ylabel('Predictions [MPG]')
```

```
plt.axis('equal')
```

```
plt.axis('square')
```

```
plt.plot([-100, 100], [-100, 100])
```

`plt.show()`



```

error = test_predictions - test_labels
plt.hist(error, bins=25)
plt.xlabel('Prediction Error [MPG]')
plt.ylabel('Count')
plt.show()

```

```

Predicted vs Actual MPG values:
Predicted: 1.33, Actual: 4.00
Predicted: 1.69, Actual: 4.00
Predicted: 4.12, Actual: 4.00
Predicted: 1.36, Actual: 4.00
Predicted: 1.94, Actual: 4.00
Predicted: 4.14, Actual: 4.00
Predicted: 8.92, Actual: 8.00
Predicted: 0.91, Actual: 4.00
Predicted: 2.69, Actual: 6.00
Predicted: 1.09, Actual: 4.00
Predicted: 6.17, Actual: 8.00
Predicted: 5.75, Actual: 6.00
Predicted: 4.98, Actual: 8.00
Predicted: 4.17, Actual: 4.00
Predicted: 2.53, Actual: 6.00
Predicted: 1.36, Actual: 4.00
Predicted: 6.02, Actual: 6.00
Predicted: 4.09, Actual: 4.00
Predicted: 1.85, Actual: 4.00
Predicted: 1.47, Actual: 4.00
Predicted: 6.06, Actual: 6.00
Predicted: 0.94, Actual: 4.00
Predicted: 0.93, Actual: 4.00
...
Predicted: 1.92, Actual: 4.00
Predicted: 4.77, Actual: 8.00
Predicted: 4.29, Actual: 8.00
Predicted: 4.06, Actual: 4.00

```